


a


Training: two label-supervised views, shared parameters



$$L = CE(f_{\theta}(x), y) + \lambda CE(f_{\theta}(x_{fg}), y), \lambda = 1$$

b

Inference: unchanged recognition graph



one ordinary-RGB image; one forward pass

no mask, donor selection, or auxiliary branch